

FinAssist AI — Internal Operations Review (DRAFT)

Quarterly notes / working copy / some figures still need checking

Overview maybe

FinAssist AI provides automated customer-support and payment-analysis tools for small financial-service companies. The team has grown quickly. Management says customer demand is strong, but this document combines notes from operations, engineering and finance and was not cleaned up before circulation.

Revenue increased from approximately €1.8m last year to €2.5m this year. Another spreadsheet says €2.42m, so Finance needs to confirm the final number. The company expects growth next year. Growth is good. Growth also increased infrastructure usage.

Things / costs / AI

Item	Earlier figure	Current note
Monthly AI requests	410,000	about 690,000
Inference spend	€18,400/month	€41,700/month
Support tickets	1,120/month	1,870/month
Average response latency	1.8 sec	3.6 sec at peak
Failed/retried payment workflows	0.7%	2.4% (engineering estimate)

Engineering notes (copied from meeting)

- Most customer-support classification requests still use the same high-capability reasoning model as complex fraud investigations.
- Someone suggested routing simple requests to a smaller model but no production policy exists yet.
- Retry logic is inconsistent between the mobile API and payment worker. In one incident, a client timeout caused a retry while the first payment was still processing.
- Idempotency keys exist in API documentation, but one legacy worker does not enforce uniqueness before contacting the payment provider.
- Monitoring dashboard reports aggregate latency but cannot currently explain cost by workload type.
- NOTE: investigate why the new model is slower on short requests. Maybe prompt length? Maybe queueing? not confirmed.

Security???

A test environment log captured complete request bodies for debugging. One engineer reported that a few test records contained customer names and partial account identifiers. The team says production logging is different, but the exact retention configuration was not included in this review. Action: verify before next customer audit.

There is also an old API token used by a scheduled reporting job. Nobody in the meeting could confirm the last rotation date. This may be harmless if the token is already scoped correctly, but it should be checked.

Customer / product notes

Customer satisfaction was reported as 91% in the product deck. Support's spreadsheet says 86%. The difference may be due to different survey windows. Duplicate-charge complaints increased after a checkout update, but the product team has not established whether the update caused the increase.

Several enterprise prospects asked for faster responses and clearer evidence about where AI-generated answers come from. One prospect asked whether current interest-rate questions are grounded in fresh sources or answered only from model memory.

Random meeting comments

We should maybe redesign the dashboard colors.

Office coffee machine replacement is delayed.

Sales wants a German-language demo next month.

Inference cost seems too high for easy questions.

Do not forget the customer audit.

Someone suggested buying larger GPU capacity, but no cost comparison was prepared.

Could route simple vs difficult tasks? Needs evidence.

Possible risks — not ranked

Risk / observation	Evidence in notes	Status
Inference cost growth	Requests +68%; monthly inference spend more than doubled	Needs analysis
Duplicate payments	Retries + inconsistent idempotency enforcement	Open
Sensitive logging	Test logs may contain customer identifiers	Verify
Latency	Peak average reported at 3.6 sec	Open
Model over-provisioning	Simple classification and complex investigations use same model	No routing policy
Freshness / grounding	Prospect questioned current-rate answers	No clear process

Numbers that may matter

If request volume remains near 690k/month and the current architecture remains unchanged, Finance expects AI infrastructure to become one of the fastest-growing operating expenses. No approved savings target exists. A rough meeting note says "maybe 25–40%" could be saved by routing, but this is NOT measured and must not be presented as fact.

Engineering recommends benchmarking at least one simple support workload and one reasoning-heavy workload before changing production routing. The team specifically wants measured accuracy, latency, token usage and estimated cost.

TODO / unresolved

1. Confirm revenue figure.
2. Audit idempotency behavior across all payment workers.
3. Verify logging retention and token rotation.
4. Benchmark simple vs reasoning workloads.
5. Determine whether fresh-data retrieval should be conditional.
6. Quantify whether smaller models preserve acceptable quality.

End of draft. This document is intentionally messy and internally inconsistent for RouteIQ document/prompt-optimization testing.